

CS 553 PA5

An Empirical Evaluation of Distributed Key/Value Storage Systems

Instructions:

- Assigned date: Monday November 13th, 2023.
- Due date: **11:59PM on Monday November 27th, 2023.**
- Maximum Points: 100%.
- This homework can be done in groups of 2 students.
- Please post your questions to BB.
- Upload your assignment on the Blackboard with the following name:
Section_LastName1_FirstName1-Section_LastName2_FirstName2.PA5.zip.
- Late submission will be penalized at 10% per day.
- Usage of online resources is allowed if citation is included. Referring to AI tools are allowed; you are responsible for the correctness of your submitted code.

1 The Problem:

In this project, you are going to evaluate various distributed key/value storage systems and compare them to existing literature.

1.1 Read paper.

Read paper on ZHT [1]. Pay special attention to Figure 8 and Figure 10, as you are going to conduct similar experiments in this project.

[1] Tonglin Li, Xiaobing Zhou, Kevin Brandstatter, Dongfang Zhao, Ke Wang, Anupam Rajendran, Zhao Zhang, Ioan Raicu. "ZHT: A Light-weight Reliable Persistent Dynamic Scalable Zero-hop Distributed Hash Table", IEEE International Parallel & Distributed Processing Symposium (IPDPS) 2013;
http://datasys.cs.iit.edu/publications/2013_IPDPS13_ZHT.pdf The 3 systems that were compared in this paper were:

- ZHT (see above paper)
- Cassandra (<http://cassandra.apache.org>)
- Memcached (<https://memcached.org>)

Please use the following data (that came from the ZHT paper [1]) to plug into your own results comparison:

System / Scale	1	2	4	8
ZHT	0.243	0.362	0.408	0.428
Cassandra	1.199	1.870	1.875	1.994
Memcached	0.122	0.324	0.272	0.278

Figure 8: Performance evaluation of ZHT and Memcached plotting latency vs. scale (1 to 64 nodes on the HECCluster); unit of measurement is milliseconds per operation.

System / Scale	1	2	4	8
ZHT	4117	5524	9813	18680
Cassandra	926	1131	2189	4322
Memcached	7385	7961	14480	40995

Figure 10: Performance evaluation of ZHT, Memcached and Casandra plotting throughput vs. scale (1 to 64 nodes on the HEC-Cluster); unit of measurement is operations per second.

1.2 Testbed: Chameleon Cloud

You are to use the Chameleon cloud (<https://www.chameleoncloud.org>) to conduct your evaluation. Evaluate your system on up to 8 instances. You must deploy Ubuntu Linux 22.04 using either “compute-skylake” or “compute-haswell” nodes, at either UC or TACC sites. Once you create a lease (up to 7 days are allowed), and start your 1 physical node, and Linux boots, you will find yourself with a physical node with 24 CPU cores, 48 hardware threads, 128GB to 192GB of memory (depending on if its Haswell or Skylake), and 250GB SSD hard drive. You will install your favorite virtualization tools (e.g. virtualbox, KVM, qemu), and use it to deploy 8 VMs with the following sizes: small.instance (3-cores, 12gb ram, 24gb disk).

1.2 Systems to Evaluate

You must evaluate the following distributed key/value stores:

- Cassandra (<http://cassandra.apache.org>)

You must also choose 1 of the following 7 distributed key/value stores to evaluate:

- MongoDB (<https://www.mongodb.org>)
- CouchDB (<http://couchdb.apache.org>)
- HBase (<http://hbase.apache.org>)
- Riak (<http://basho.com/products/#riak>)
- Redis (<http://redis.io>)
- HyperTable (<http://www.hypertable.com>)
- Oracle NoSQL Database (<https://www.oracle.com/database/technologies/related/nosql.html>)

You are to compare these 2 systems to the data presented in paper [1].

1.3 Compare and contrast your 2 chosen systems to ZHT

Write a paragraph about each of your 2 systems you chose, as well as about ZHT summarizing what the systems are, and what the key features are. Follow this writeup with a comparison that discusses the similarities and differences between your chosen 2 systems to ZHT. Create a table of features that clearly shows how these systems are different. This section should be about 2 pages long (+/- 0.5 pages). Too short, and you will lose points for not enough detail. Too long, and you will also lose points for being verbose.

1.4 Evaluation Scale and Metrics

Read this paper on ZHT [1]. Particularly, Figure 8 and Figure 10 are very important. You will have to conduct enough experiments to generate 2 figures, 1 for latency and 1 for throughput, with each data point representing the average across all 3 operations (insert, lookup, and remove).

On each instance/node, a client-server pair is deployed. Test workload is a set of key-value pairs where the key is 10 bytes and value is 90 bytes. Clients sequentially send all the key-value pairs through a client API for insert, then lookup, and then remove. Your keys should be randomly generated, which will produce an All-to-All communication pattern, with the same number of servers and clients.

The metrics you will measure, and report are:

- **Latency:** Latency presents the time per operation (insert/lookup/remove) taken from a request to be submitted from a client to a response to be received by the client, measured in milliseconds (ms). Note that the latency consists of round-trip network communication, system processing, and storage access time.
- **Throughput:** The number of operations (insert/lookup/remove) the system can handle over some period of time, measured in Ops/s

Make the necessary plots to visualize your data. Explain why your results make sense; what explicit things did you do to verify that your performance as you expected? If there are differences in your results compared to those published in [1], discuss why these differences might make sense. What changes would you have to do to your experiments to validate your theories about the differences in your results?

2 Where you will submit:

When you have finished implementing the complete assignment as described above, you should submit your solution to the blackboard. The timestamp on the BB submission will be used to determine if the submission is on time. Please put all your homework content into one .zip file and upload it to the blackboard.

The name of .zip should follow this format: "**Section_LastName1_FirstName1-Section_LastName2_FirstName2.PA5.zip**".

Submit step #1: Put all your files and documents for submission in a zip file.

Submit step #2: Name the file correctly and submit it to the blackboard.

Grades for late programs will be lowered 10% per day late.

3 What you will submit & Grading:

Each program must work correctly and be detailed in-line documented. You should hand in:

1. **Source Code:** You must hand in all your source code of any scripts you used to automate your performance evaluation. If you wrote any programs specific to each system to perform the evaluation, include these evaluation drivers.
2. **Report:** You must write a project report to cover all the requirements in this assignment; you will be scored on completeness, correctness, organization, and visual appeal.

Grades for late programs will be lowered 10% per day late.